

```
In [1]: from pandas import read_excel, DataFrame, crosstab
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib.colors import Normalize
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, LabelEncoder, StandardScaler
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from imblearn.pipeline import Pipeline as ImbPipeline
from imblearn.over_sampling import RandomOverSampler, SMOTE, SVMSMOTE, BorderlineSMOTE
import joblib
import json
from os.path import exists
```

```
In [2]: df = read_excel("IIMK_DSAI_W11_Required Assignment 11.1_Data Set.xlsm")
```

```
In [3]: print(df.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
In [4]: print(df.tail())
```

	PassengerId	Survived	Pclass	Name \
886	887	0	2	Montvila, Rev. Juozas
887	888	1	1	Graham, Miss. Margaret Edith
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"
889	890	1	1	Behr, Mr. Karl Howell
890	891	0	3	Dooley, Mr. Patrick

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
886	male	27.0	0	0	211536	13.00	NaN	S
887	female	19.0	0	0	112053	30.00	B42	S
888	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	male	26.0	0	0	111369	30.00	C148	C
890	male	32.0	0	0	370376	7.75	NaN	Q

```
In [5]: # Check if there are any duplicate values in dataset
print(sum(df.duplicated()))
```

0

```
In [6]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age         714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None
```

```
In [7]: na_counts = DataFrame(df.isna().sum(),columns=["NA Counts"]).reset_index()
na_counts = na_counts.rename(columns={'index': 'Column Name'})
print(na_counts)
```

	Column Name	NA Counts
0	PassengerId	0
1	Survived	0
2	Pclass	0
3	Name	0
4	Sex	0
5	Age	177
6	SibSp	0
7	Parch	0
8	Ticket	0
9	Fare	0
10	Cabin	687
11	Embarked	2

In [8]: `df.describe()`

Out[8]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [9]: `# Printing unique values in datasets columns`
`for column in ["Survived", "Pclass", "Sex", "SibSp", "Parch", "Cabin", "Embarked"]:`
 `print(f"{column}:{df[column].unique()}\n")`

Survived:[0 1]

Pclass:[3 1 2]

Sex:['male' 'female']

SibSp:[1 0 3 4 2 5 8]

Parch:[0 1 2 5 3 4 6]

Cabin:[nan 'C85' 'C123' 'E46' 'G6' 'C103' 'D56' 'A6' 'C23 C25 C27' 'B78' 'D33'
'B30' 'C52' 'B28' 'C83' 'F33' 'F G73' 'E31' 'A5' 'D10 D12' 'D26' 'C110'
'B58 B60' 'E101' 'F E69' 'D47' 'B86' 'F2' 'C2' 'E33' 'B19' 'A7' 'C49'
'F4' 'A32' 'B4' 'B80' 'A31' 'D36' 'D15' 'C93' 'C78' 'D35' 'C87' 'B77'
'E67' 'B94' 'C125' 'C99' 'C118' 'D7' 'A19' 'B49' 'D' 'C22 C26' 'C106'
'C65' 'E36' 'C54' 'B57 B59 B63 B66' 'C7' 'E34' 'C32' 'B18' 'C124' 'C91'
'E40' 'T' 'C128' 'D37' 'B35' 'E50' 'C82' 'B96 B98' 'E10' 'E44' 'A34'
'C104' 'C111' 'C92' 'E38' 'D21' 'E12' 'E63' 'A14' 'B37' 'C30' 'D20' 'B79'
'E25' 'D46' 'B73' 'C95' 'B38' 'B39' 'B22' 'C86' 'C70' 'A16' 'C101' 'C68'
'A10' 'E68' 'B41' 'A20' 'D19' 'D50' 'D9' 'A23' 'B50' 'A26' 'D48' 'E58'
'C126' 'B71' 'B51 B53 B55' 'D49' 'B5' 'B20' 'F G63' 'C62 C64' 'E24' 'C90'
'C45' 'E8' 'B101' 'D45' 'C46' 'D30' 'E121' 'D11' 'E77' 'F38' 'B3' 'D6'
'B82 B84' 'D17' 'A36' 'B102' 'B69' 'E49' 'C47' 'D28' 'E17' 'A24' 'C50'
'B42' 'C148']

Embarked:['S' 'C' 'Q' nan]

```
In [10]: #Printing unique values in dataset columns
for column in ["Survived", "Pclass", "Sex", "SibSp", "Parch", "Cabin", "Embarked"]:
    print(f"{df[column].value_counts()}{df[column].value_counts(normalize=True)}\n")
```

Survived
0 549
1 342
Name: count, dtype: int64
Survived
0 0.616162
1 0.383838
Name: proportion, dtype: float64

Pclass
3 491
1 216
2 184
Name: count, dtype: int64
Pclass
3 0.551066
1 0.242424
2 0.206510
Name: proportion, dtype: float64

Sex
male 577
female 314
Name: count, dtype: int64
Sex
male 0.647587
female 0.352413
Name: proportion, dtype: float64

SibSp
0 608
1 209
2 28
4 18
3 16
8 7
5 5
Name: count, dtype: int64
SibSp
0 0.682379
1 0.234568
2 0.031425
4 0.020202
3 0.017957
8 0.007856
5 0.005612
Name: proportion, dtype: float64

Parch
0 678
1 118
2 80
5 5
3 5
4 4
6 1
Name: count, dtype: int64
Parch
0 0.760943
1 0.132435
2 0.089787

```

5    0.005612
3    0.005612
4    0.004489
6    0.001122
Name: proportion, dtype: float64

```

```

Cabin
B96 B98      4
G6          4
C23 C25 C27  4
C22 C26      3
F33          3
..
E34          1
C7           1
C54          1
E36          1
C148         1

```

Name: count, Length: 147, dtype: int64Cabin

```

B96 B98      0.019608
G6          0.019608
C23 C25 C27  0.019608
C22 C26      0.014706
F33          0.014706

```

```

...
E34          0.004902
C7           0.004902
C54          0.004902
E36          0.004902
C148         0.004902

```

Name: proportion, Length: 147, dtype: float64

Embarked

```

S    644
C    168
Q     77

```

Name: count, dtype: int64Embarked

```

S    0.724409
C    0.188976
Q    0.086614

```

Name: proportion, dtype: float64

```

In [11]: # Replacing NA in Age with the mean age
mean_Age = df['Age'].mean()
df['Age'] = df['Age'].fillna(mean_Age)

```

```

In [12]: # As there are several NA in Cabin, for better analysis making another column is_Ca
df['is_Cabin'] = np.where(df['Cabin'].isna(), 0, 1)
df = df.drop('Cabin', axis = 1)

```

```

In [13]: # Replacing NA in Embarked with the mode
mode_Embarked = df['Embarked'].mode()
df['Embarked'] = df['Embarked'].fillna(mode_Embarked[0])

```

```
In [14]: # Now checking for NA
na_counts = DataFrame(df.isna().sum(),columns=["NA Counts"]).reset_index()
na_counts = na_counts.rename(columns={'index': 'Column Name'})
print(na_counts)
```

	Column Name	NA Counts
0	PassengerId	0
1	Survived	0
2	Pclass	0
3	Name	0
4	Sex	0
5	Age	0
6	SibSp	0
7	Parch	0
8	Ticket	0
9	Fare	0
10	Embarked	0
11	is_Cabin	0

```
In [15]: # Printing unique values in datasets columns
for column in ["Survived","Pclass","Sex","SibSp","Parch","is_Cabin","Embarked"]:
    print(f"{column}:{df[column].unique()}\n")
```

Survived:[0 1]

Pclass:[3 1 2]

Sex:['male' 'female']

SibSp:[1 0 3 4 2 5 8]

Parch:[0 1 2 5 3 4 6]

is_Cabin:[0 1]

Embarked:['S' 'C' 'Q']

```
In [16]: #Printing unique values in dataset columns
for column in ["Survived","Pclass","Sex","SibSp","Parch","is_Cabin","Embarked"]:
    print(f"{df[column].value_counts()}{df[column].value_counts(normalize=True)}\n")
```

Survived
0 549
1 342
Name: count, dtype: int64
Survived
0 0.616162
1 0.383838
Name: proportion, dtype: float64

Pclass
3 491
1 216
2 184
Name: count, dtype: int64
Pclass
3 0.551066
1 0.242424
2 0.206510
Name: proportion, dtype: float64

Sex
male 577
female 314
Name: count, dtype: int64
Sex
male 0.647587
female 0.352413
Name: proportion, dtype: float64

SibSp
0 608
1 209
2 28
4 18
3 16
8 7
5 5
Name: count, dtype: int64
SibSp
0 0.682379
1 0.234568
2 0.031425
4 0.020202
3 0.017957
8 0.007856
5 0.005612
Name: proportion, dtype: float64

Parch
0 678
1 118
2 80
5 5
3 5
4 4
6 1
Name: count, dtype: int64
Parch
0 0.760943
1 0.132435
2 0.089787


```
5    0.005612
3    0.005612
4    0.004489
6    0.001122
Name: proportion, dtype: float64
```

```
is_Cabin
0    687
1    204
Name: count, dtype: int64is_Cabin
0    0.771044
1    0.228956
Name: proportion, dtype: float64
```


```
Embarked
S    646
C    168
Q     77
Name: count, dtype: int64Embarked
S    0.725028
C    0.188552
Q    0.086420
Name: proportion, dtype: float64
```

Exploratory Data Analysis

As we know that females, elderly and children were preferred to be saved via lifeboats let us check the survival on basis of these two parameters. Then let us analyze it in terms of Pclass to check if there was a preference to save 1st class passengers.

```
In [17]: # Create a heatmap directly from the DataFrame
sns.heatmap(crosstab(df['Survived'], df['Sex']), cmap='Blues', annot=True) # Adjust

# Customize the plot
plt.xlabel('Gender')
plt.ylabel('Survived')
plt.title('Heatmap - Survived vs Gender')
plt.show()
```




```
In [18]: # Create the violin plot
sns.violinplot(x = "Survived",
               y = "Age",
               data = df,
               hue = "Survived",
               palette = "viridis")

# Customize the plot
plt.xlabel("Survived")
plt.ylabel("Age")
plt.title("Violin Plot of Survived vs Age")
plt.xticks(rotation=45)
plt.show()
```



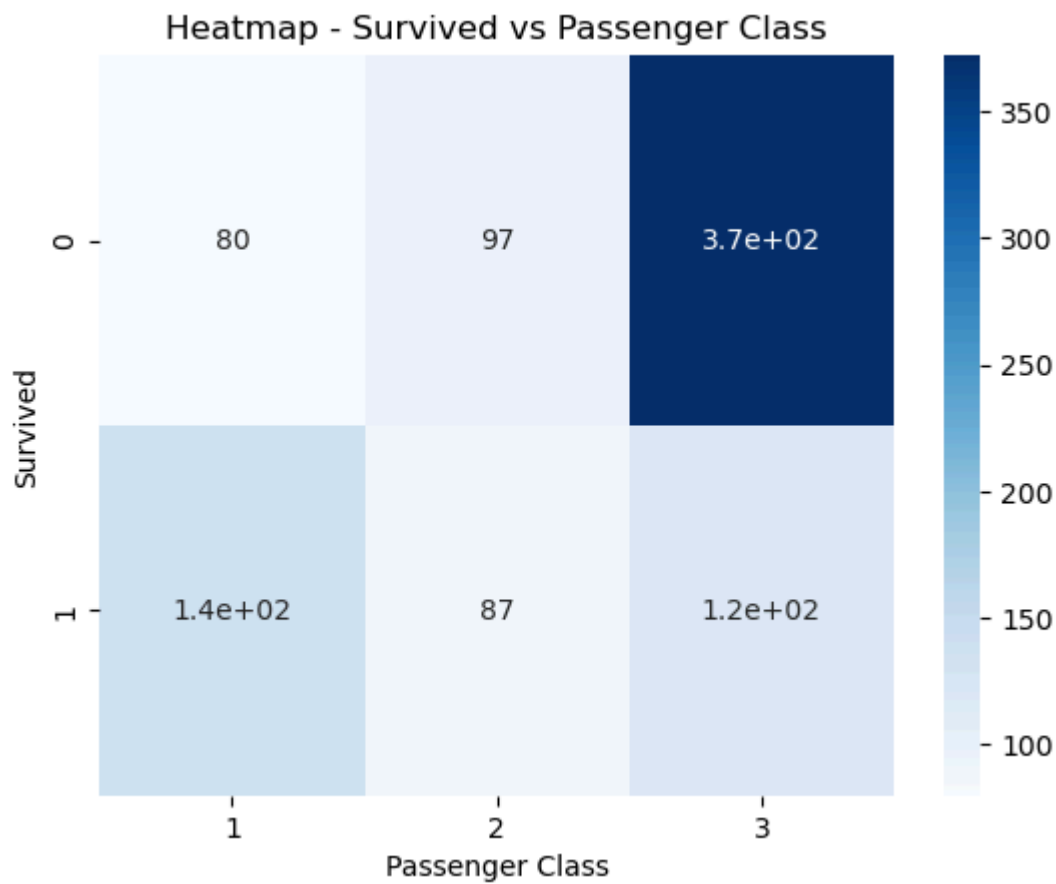
```
In [19]: # Create the violin plot
sns.violinplot(x = "Survived",
               y = "Age",
               data = df,
               hue = "Sex",
               palette = "viridis")

# Customize the plot
plt.xlabel("Survived")
plt.ylabel("Age")
plt.title("Violin Plot of Survived vs Age and Gender")
plt.xticks(rotation=45)
plt.show()
```




```
In [20]: # Create a heatmap directly from the DataFrame
sns.heatmap(crosstab(df['Survived'], df['Pclass']), cmap='Blues', annot=True) # Ad

# Customize the plot
plt.xlabel('Passenger Class')
plt.ylabel('Survived')
plt.title('Heatmap - Survived vs Passenger Class')
plt.show()
```



```
In [21]: # Create a heatmap directly from the DataFrame
sns.heatmap(crosstab(df['Survived'], df['is_Cabin']), cmap='Blues', annot=True) #

# Customize the plot
plt.xlabel('Cabin Passenger')
plt.ylabel('Survived')
plt.title('Heatmap - Survived vs Cabin Passenger')
plt.show()
```



We see a clear indication that female passengers are more likely to survive. Children and the Elderly were also more likely to survive especially in males. Passengers in 1st Class are more likely to survive than passengers in 2nd Class and 3rd Class. Passengers in Cabin are also more likely to survive than passengers not having cabin tickets.

Splitting the Training and Testing Data Set

```
In [22]: # The only features we will use for further modelling - SibSp, Pclass, Sex, Age, Pa
X,y = df.drop(["PassengerId", "Name", "Ticket", "Survived"],axis=1), df["Survived"]

# Split data into training and testing sets (default test_size=0.25) # Through tria
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st
```

```
In [23]: print("Training Data Shape")
print(X_train.shape)
print("Testing Data Shape")
print(X_test.shape)
```

```
Training Data Shape
(668, 8)
Testing Data Shape
(223, 8)
```

```
In [24]: # Define column names
ordinal_cols = ['Sex']
```

```

onehot_cols = ['Embarked']
numerical_cols = [col for col in X.columns if col not in ordinal_cols + onehot_cols]

# Define the ColumnTransformer
preprocessor = ColumnTransformer(
    transformers=[
        ('ordinal', OrdinalEncoder(), ordinal_cols),
        ('onehot', OneHotEncoder(), onehot_cols),
        ('num', 'passthrough', numerical_cols)
    ]
)

```


```

In [25]: transformed_df = preprocessor.fit_transform(df.drop(["PassengerId", "Name", "Ticket"]))
transformed_df = DataFrame(transformed_df, columns=['Sex', 'Embarked_C', 'Embarked_S',
transformed_df["Survived"] = df["Survived"].astype(float)
transformed_df["Embarked_S"] = transformed_df.apply(lambda row: 1.0 if row['Embarke

corr = transformed_df.corr()
norm = Normalize(vmin=-0.54, vmax=0.54)
plt.figure(figsize=(10, 7.5))
sns.heatmap(corr["Survived"].to_frame().sort_values(by="Survived", key=lambda x: x.
            annot=True, cmap = "coolwarm", norm = norm)

```

Out[25]: <Axes: >



As per the above heatmap we can analyze the feature importance and correlation with the target variable Survived. By this analysis, we find the maximum importance is of gender

followed by Passenger Class, whether the passenger has a cabin ticket and fare. The other features have relatively lesser importance due to its low correlation with the target variable.

SVM

```
In [26]: # If model already run from the existing model or else define the model
if exists('svm_model.joblib'):
    print("Loading from file")
    svm_loaded = joblib.load('svm_model.joblib')
    y_pred_svm = svm_loaded.predict(X_test)
    y_pred_proba_svm = svm_loaded.predict_proba(X_test)
    with open('svm_results.json', 'r') as json_file:
        results = json.load(json_file)
        optimal_params = results['optimal_params']
        optimal_accuracy = results['optimal_accuracy']
        print(f'Best parameters: {optimal_params}')
        print(f'Best cross-validation accuracy: {optimal_accuracy:.5f}')
else:
    # Define the pipeline
    pipeline = ImbPipeline(steps=[
        ('preprocessor', preprocessor),
        ('scaler', StandardScaler()), # Apply scaling to all columns after preproc
        ('smote_svm', SVMSMOTE(random_state=21)),
        ('svm', SVC(probability=True))
    ])

    # Define the parameter grid for GridSearchCV
    param_grid = {
        'svm_C': np.arange(0.25, 1.8, 0.25),
        'svm_kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
    }

    # Define the GridSearchCV
    grid_search = GridSearchCV(pipeline, param_grid, cv=5, njobs = -1)

    # Perform GridSearchCV to find the best parameters and fit the model
    grid_search.fit(X_train, y_train)

    # Update the pipeline with the best estimator
    optimal_estimator = grid_search.best_estimator_

    y_pred_svm = optimal_estimator.predict(X_test)
    y_pred_proba_svm = optimal_estimator.predict_proba(X_test)

    # Print the best parameters and cross-validation accuracy
    print(f'Best parameters: {(optimal_params:= grid_search.best_params_)}')
    print(f'Best cross-validation accuracy: {(optimal_accuracy:= grid_search.best_s

    # Writing it in json file
    results = {'optimal_params': optimal_params, 'optimal_accuracy': optimal_accura
    with open('svm_results.json', 'w') as json_file:
        json.dump(results, json_file)
```



```
# Save the pipeline
joblib.dump(optimal_estimator, 'svm_model.joblib')
print("Model trained and saved to disk.")
```

Loading from file

Best parameters: {'svm__C': 1.5, 'svm__kernel': 'rbf'}

Best cross-validation accuracy: 0.82781

Model Evaluation

```
In [27]: accuracy_svm = accuracy_score(y_test, y_pred_svm)
print(f"The accuracy of SVM model is {accuracy_svm:5f}")
```

The accuracy of SVM model is 0.825112

```
In [28]: print(confusion_matrix(y_test, y_pred_svm))
```

```
[[122  11]
 [ 28  62]]
```

```
In [29]: print(classification_report(y_test,y_pred_svm))
```

	precision	recall	f1-score	support
0	0.81	0.92	0.86	133
1	0.85	0.69	0.76	90
accuracy			0.83	223
macro avg	0.83	0.80	0.81	223
weighted avg	0.83	0.83	0.82	223


```
In [30]: f1_svm = f1_score(y_test,y_pred_svm)
print(f"The f1 score of SVM model is {f1_svm:5f}")
```

The f1 score of SVM model is 0.760736

```
In [31]: roc_auc_svm = roc_auc_score(y_test, y_pred_proba_svm[:,1])
print(f"The ROC-AUC score of SVM model is {roc_auc_svm:5f}")
```

The ROC-AUC score of SVM model is 0.854511

```
In [32]: fpr_svm, tpr_svm, _ = roc_curve(y_test, y_pred_proba_svm[:, 1])
plt.plot(fpr_svm, tpr_svm, label='SVM (AUC = {:.5f})'.format(roc_auc_svm))
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for SVM')
plt.text(0.6, 0.2, f'AUC = {roc_auc_svm:.5f}', fontsize=12, color='navy')
plt.show()
```



Gaussian Naive Bayes

```
In [33]: # If model already run from the existing model or else define the model
if exists('gnb_model.joblib'):
    print("Loading from file")
    gnb_loaded = joblib.load('gnb_model.joblib')
    y_pred_gnb = gnb_loaded.predict(X_test)
    y_pred_proba_gnb = gnb_loaded.predict_proba(X_test)
else:
    # Define the pipeline
    pipeline = ImbPipeline(steps=[
        ('preprocessor', preprocessor),
        ('smote', ADASYN(random_state=21)),
        ('scaler', StandardScaler()), # Apply scaling to all columns after preproc
        ('gnb', GaussianNB())
    ])

    pipeline.fit(X_train, y_train)

    y_pred_gnb = pipeline.predict(X_test)
    y_pred_proba_gnb = pipeline.predict_proba(X_test)

    # Save the pipeline
    joblib.dump(pipeline, 'gnb_model.joblib')
    print("Model trained and saved to disk.")
```

Loading from file

Model Evaluation

```
In [34]: accuracy_gnb = accuracy_score(y_test, y_pred_gnb)
print(f"The accuracy of Gaussian Naive Bayes model is {accuracy_gnb:5f}")
```

The accuracy of Gaussian Naive Bayes model is 0.807175

```
In [35]: print(confusion_matrix(y_test, y_pred_gnb))

[[105  28]
 [ 15  75]]
```

```
In [36]: print(classification_report(y_test,y_pred_gnb))
```

	precision	recall	f1-score	support
0	0.88	0.79	0.83	133
1	0.73	0.83	0.78	90
accuracy			0.81	223
macro avg	0.80	0.81	0.80	223
weighted avg	0.82	0.81	0.81	223

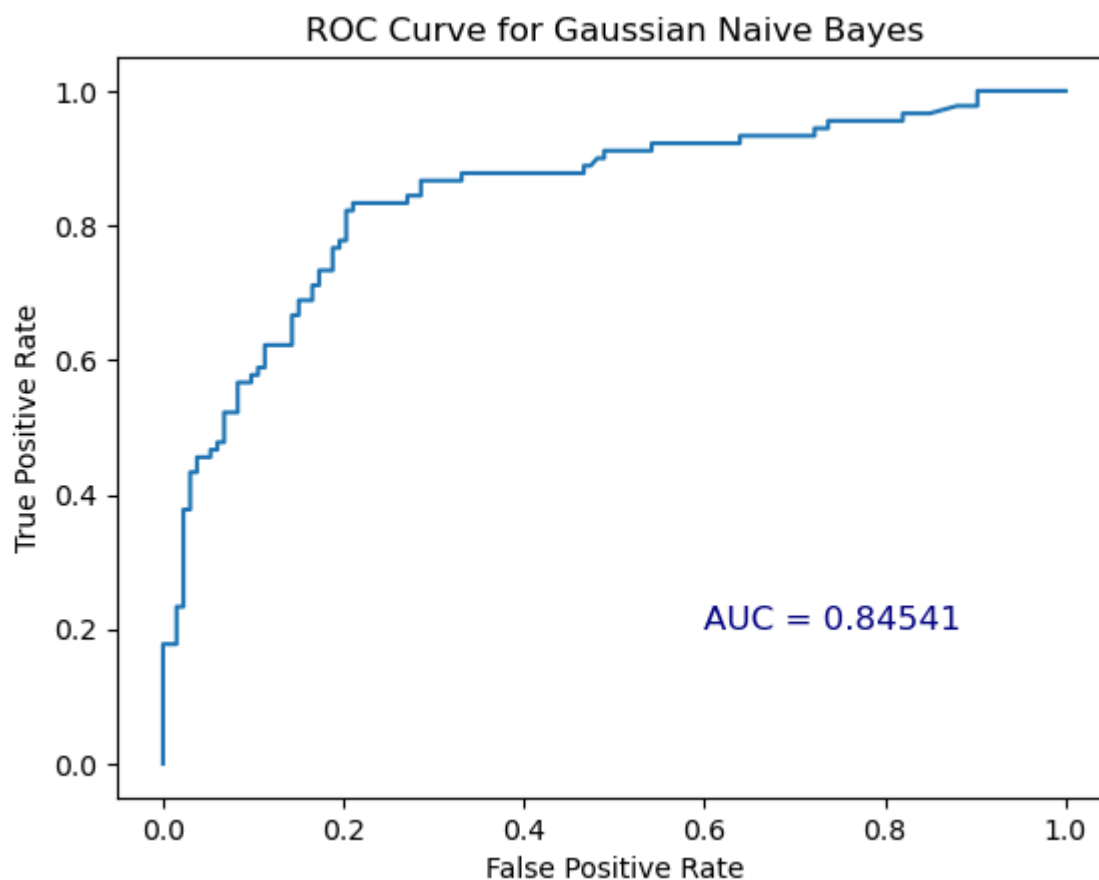
```
In [37]: f1_gnb = f1_score(y_test,y_pred_gnb)
print(f"The f1 score of Gaussian Naive Bayes model is {f1_gnb:5f}")
```

The f1 score of Gaussian Naive Bayes model is 0.777202

```
In [38]: roc_auc_gnb = roc_auc_score(y_test, y_pred_proba_gnb[:,1])
print(f"The ROC-AUC score of GNB model is {roc_auc_gnb:5f}")
```

The ROC-AUC score of GNB model is 0.845405

```
In [39]: fpr_gnb, tpr_gnb, _ = roc_curve(y_test, y_pred_proba_gnb[:, 1])
plt.plot(fpr_gnb, tpr_gnb, label='GNB (AUC = {:.5f})'.format(roc_auc_gnb))
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Gaussian Naive Bayes')
plt.text(0.6, 0.2, f'AUC = {roc_auc_gnb:.5f}', fontsize=12, color='navy')
plt.show()
```



Comparison of Models

```
In [40]: metrics = DataFrame({"SVM": [accuracy_svm, f1_svm, roc_auc_svm],
                             "GNB": [accuracy_gnb, f1_gnb, roc_auc_gnb]},
                             index = ["Accuracy", "F1 Score", "ROC AUC Score"])
metrics = metrics.rename_axis('Metrics')
display(metrics)
```

	SVM	GNB
Metrics		
Accuracy	0.825112	0.807175
F1 Score	0.760736	0.777202
ROC AUC Score	0.854511	0.845405

```
In [41]: plt.plot(fpr_svm, tpr_svm, label='SVM (AUC = {:.5f})'.format(roc_auc_svm), color = 'blue')
plt.plot(fpr_gnb, tpr_gnb, label='GNB (AUC = {:.5f})'.format(roc_auc_gnb), color = 'red')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Gaussian Naive Bayes')
plt.legend(loc='lower right')
plt.show()
```

ROC Curve for Gaussian Naive Bayes

